

```
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Dense, LeakyReLU, BatchNormalization, Reshape, Flatten, Input
from tensorflow.keras.optimizers import Adam
```

```
(X_train, _), (_, _) = mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step
```

```
X_train = X_train / 127.5 - 1.0
X_train = np.expand_dims(X_train, axis=3)
```

```
def build_generator(latent_dim):
    model = Sequential()
    model.add(Dense(128 * 7 * 7, input_dim=latent_dim))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Reshape((7, 7, 128)))
    model.add(BatchNormalization())
    model.add(Flatten())
    model.add(Dense(28 * 28, activation='tanh'))
    model.add(Reshape((28, 28, 1)))
    noise = Input(shape=(latent_dim,))
    img = model(noise)
    return Model(noise, img)
```

```
def build_discriminator(img_shape):
    model = Sequential()
    model.add(Flatten(input_shape=img_shape))
    model.add(Dense(128))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dense(1, activation='sigmoid'))
    img = Input(shape=img_shape)
    validity = model(img)
    return Model(img, validity)
```

```
img_shape = X_train[0].shape
discriminator = build_discriminator(img_shape)
discriminator.compile(loss='binary_crossentropy', optimizer=Adam(lr=0.0002, beta_1=0.5), metrics=['accuracy'])
```

```
WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.leg
```

```
latent_dim = 100
generator = build_generator(latent_dim)
```


```
z = Input(shape=(latent_dim,))
img = generator(z)
discriminator.trainable = False
validity = discriminator(img)
combined = Model(z, validity)
combined.compile(loss='binary_crossentropy', optimizer=Adam(lr=0.0002, beta_1=0.5))
```

```
WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.leg
```

```
epochs = 10000
batch_size = 128
sample_interval = 1000
```

```
for epoch in range(epochs):
    # Train discriminator
    idx = np.random.randint(0, X_train.shape[0], batch_size)
    real_imgs = X_train[idx]
    noise = np.random.normal(0, 1, (batch_size, latent_dim))
```

```
noise = np.random.normal(0, 1, (batch_size, latent_dim))
fake_imgs = generator.predict(noise)
d_loss_real = discriminator.train_on_batch(real_imgs, np.ones((batch_size, 1)))
d_loss_fake = discriminator.train_on_batch(fake_imgs, np.zeros((batch_size, 1)))
d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
```

 Streaming output truncated to the last 5000 lines.

```
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 24ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 26ms/step
4/4 [=====] - 0s 27ms/step
4/4 [=====] - 0s 33ms/step
4/4 [=====] - 0s 22ms/step
4/4 [=====] - 0s 22ms/step
4/4 [=====] - 0s 30ms/step
4/4 [=====] - 0s 33ms/step
4/4 [=====] - 0s 29ms/step
4/4 [=====] - 0s 22ms/step
4/4 [=====] - 0s 30ms/step
4/4 [=====] - 0s 23ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 22ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 21ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 22ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 21ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 22ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 21ms/step
4/4 [=====] - 0s 21ms/step
4/4 [=====] - 0s 21ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 21ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 26ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step
```

```
noise = np.random.normal(0, 1, (batch_size, latent_dim))
g_loss = combined.train_on_batch(noise, np.ones((batch_size, 1)))
```

```
if epoch % sample_interval == 0:
    print(f"{epoch} [D loss: {d_loss[0]}, acc.: {100 * d_loss[1]}] [G loss: {g_loss}]")
```

